

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/> (<https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from tqdm import tqdm
from bs4 import BeautifulSoup

import re
import string
import nltk
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.stem import PorterStemmer
from nltk.stem import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer, TfidfTransformer
from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score, auc, roc_curve, classification_report, precision_score, recall_score, f1_score, hamming_loss

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import TruncatedSVD
from prettytable import PrettyTable

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```
In [2]: # using SQLite Table to read data.
con = sqlite3.connect('D:\Study_materials\Applied_AI\Assignments\database.sqlite')

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)

Number of data points in our data (525814, 10)
```

```
In [3]: sample_data = filtered_data.head(100000)
```

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [4]: #Sorting data according to ProductId in ascending order
sorted_data=sample_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

```
In [5]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

```
Out[5]: (87775, 10)
```

```
In [6]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(sample_data['Id'].size*1.0)*100
```

```
Out[6]: 87.775
```

```
In [7]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [8]: #Before starting the next phase of preprocessing Lets see the number of entries Left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(87773, 10)

Out[8]: 1    73592
0     14181
Name: Score, dtype: int64
```

```
In [9]: final.head()

Out[9]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
22620	24750	2734888454	A13ISQV0U9GZIC	Sandikaye	1	1	0	1192060800	made in china	My dogs loves this chicken but its a product f...
22621	24751	2734888454	A1C298ITT645B6	Hugh G. Pritchard	0	0	1	1195948800	Dog Lover Delites	Our dogs just love them. I saw them in a pet ...
70677	76870	B00002N8SM	A19Q006CSFT011	Arielle	0	0	0	1288396800	only one fruitfly stuck	I had an infestation of fruitflies, they were ...
70676	76869	B00002N8SM	A1FYH4S02BW7FN	wonderer	0	0	0	1290038400	Doesn't work!! Don't waste your money!!	Worst product I have gotten in long time. Woul...
70675	76868	B00002N8SM	AUE8TB5VHS6ZV	eyeofthestorm	0	0	0	1306972800	A big rip off	I wish I'd read the reviews before making this...

[3] Preprocessing

```
In [10]: import re
i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

4
I wish I'd read the reviews before making this purchase. It's basically a cardsotck box that is sticky all over the OUTSIDE. Thos e pink-ish things that look like entrances "into" the trap? They're just pictures. There *is no* inside of the trap. All the flie s will be stuck to the OUTSIDE. It's basically fly paper, just horribly, horribly HORRIBLY overpriced.

Do yourself a f avor and just get fly paper or fly strips. Same yuck factor, but much cheaper.

```
In [11]: stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext

def cleanpunc(sentence): #function to clean the word of any punctuation or special characters
    cleaned = re.sub(r'[!|!|\\"|"]|#',r'',sentence)
    cleaned = re.sub(r'[\.\,|,|(|\|/|)]',r'',cleaned)
    return cleaned

print(stop)
print('*****')
print(sno.stem('tasty'))

{"doesn't", 'before', 'into', 'being', 'didn't', 'off', "won't", 'more', "mustn't", 'had', 're', 'an', 'wouldn', "didn't", 'throug
h', 'haven', 'mustn', 'this', 'own', "couldn't", "wouldn't", 'd', 'the', 'she', "aren't", 'whom', 'ourselves', 'doesn', 'out', 'h
ers', 'now', 'hasn', 'of', "shan't", 'if', 'than', 'having', 'ours', 't', 'how', 'will', 'what', 'am', 'until', 'between', 'it',
'be', "you're", 'most', 'theirs', 'yourself', 'again', "it's", 'do', 'doing', 'couldn', 'does', 'but', 'll', 'in', 'few', 'such',
'who', 'himself', "don't", "she's", 'any', 'its', "hasn't", 'against', 'aren', 'were', "you'll", 'can', 'because', 'about', 'shou
ld', 'isn', 'her', 'by', 'themselves', 'for', 'have', 'on', 'a', 'each', 'm', 'not', 'further', 'they', 'me', 'is', 'only', 'bot
h', 'and', 'with', 'herself', "haven't", 'don', 'their', 'we', 'did', 'under', 'nor', 've', 'myself', 'here', 'as', "should've",
's', 'after', 'i', 'all', 'yourselves', 'over', 'been', 'him', 'below', 'your', 'while', "wasn't", 'our', 'ain', 'once', 'weren',
'has', 'them', "you've", 'at', 'hadn', 'where', 'some', 'was', 'when', 'wasn', 'o', 'just', 'from', "mightn't", 'yours', 'why',
'won', 'his', "shouldn't", "isn't", "that'll", "you'd", 'to', 'are', 'you', 'mightn', 'needn', "needn't", 'he', "hadn't", 'or',
'up', 'these', 'my', 'then', 'other', 'shouldn', 'there', 'which', 'ma', 'those', 'above', 'during', 'itself', 'no', 'down', 'ver
y', 'shan', 'that', 'y', 'so', 'same', 'too', "weren't"}
*****
tasti
```

```
Size of X_train and y_train: (61441,) (61441,)
Size of X_test and y_test: (26332,) (26332,)
Size of X_tr and y_tr: (43008,) (43008,)
Size of X_cv and y_cv: (18433,) (18433,)
```

RandomForest Classifier

```
In [29]: from sklearn.ensemble import RandomForestClassifier

def RF_Classifier(X_train,X_cv,y_train,y_cv):
    pred_train = []
    pred_cv = []
    t_depth = [2,3,5,8,10,20]
    estimators = [100,200,300,400,500]
    for i in t_depth:
        for j in estimators:
            clf = RandomForestClassifier(n_estimators = j, max_depth = i, n_jobs = -1, class_weight='balanced')
            clf.fit(X_train,y_train)
            prob_train = clf.predict_proba(X_train)[:,:1]
            prob_cv = clf.predict_proba(X_cv)[:,:1]
            auc_score_train = roc_auc_score(y_train,prob_train)
            auc_score_cv = roc_auc_score(y_cv,prob_cv)
            pred_train.append(auc_score_train)
            pred_cv.append(auc_score_cv)
    cmap=sns.light_palette("green")
    # representing heat map for auc score
    print("-"*40, "AUC Score for train data", "-"*40)
    pred_train = np.array(pred_train)
    pred_train = pred_train.reshape(len(t_depth),len(estimators))
    plt.figure(figsize=(10,5))
    sns.heatmap(pred_train,annot=True, cmap=cmap, fmt=".3f", xticklabels=estimators,yticklabels=t_depth)
    plt.xlabel('Estimator')
    plt.ylabel('Depth')
    plt.show()
    print("-"*40, "AUC Score for CV data", "-"*40)
    pred_cv = np.array(pred_cv)
    pred_cv = pred_cv.reshape(len(t_depth),len(estimators))
    plt.figure(figsize=(10,5))
    sns.heatmap(pred_cv, annot=True, cmap=cmap, fmt=".3f", xticklabels=estimators, yticklabels=t_depth)
    plt.xlabel('Estimator')
    plt.ylabel('Depth')
    plt.show()
```

Testing model:

```
In [33]: import scikitplot.metrics as skplt

def testing(X_train,y_train,X_test,y_test,optimal_depth,optimal_estimator):
    clf = RandomForestClassifier(n_estimators = optimal_estimator, max_depth = optimal_depth,class_weight='balanced')
    clf.fit(X_train,y_train)
    prob_train = clf.predict_proba(X_train)[:,:1]
    prob_test = clf.predict_proba(X_test)[:,:1]

    print("AUC Score for train data",roc_auc_score(y_train,prob_train))
    print("AUC Score for test data",roc_auc_score(y_test,prob_test))
    # calculate roc curve
    fpr_train, tpr_train, threshold_tr = roc_curve(y_train,prob_train)
    fpr_test, tpr_test, threshold_te = roc_curve(y_test,prob_test)

    # plot the roc curve for the model

    plt.plot([0, 1], [0, 1], linestyle='--')
    plt.plot(fpr_train, tpr_train, marker='.',color= 'r',label='Train Data')
    plt.plot(fpr_test, tpr_test, marker='.',color= 'b',label='Test Data')
    plt.title("Line Plot of ROC Curve on Train Data and Test Data")
    plt.legend(loc='upper left')
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()

    #plot confusion matrix

    prediction_train=clf.predict(X_train)
    prediction_test=clf.predict(X_test)

    print("macro f1 score for train data :",f1_score(y_train, prediction_train, average = 'macro'))
    print("macro f1 score for test data :",f1_score(y_test, prediction_test, average = 'macro'))
    print("micro f1 score for train data:",f1_score(y_train, prediction_train, average = 'micro'))
    print("micro f1 score for test data:",f1_score(y_test, prediction_test, average = 'micro'))
    print("hamming loss for train data:",hamming_loss(y_train,prediction_train))
    print("hamming loss for test data:",hamming_loss(y_test,prediction_test))
    print("Precision recall report for train data:\n",classification_report(y_train, prediction_train))
    print("Precision recall report for test data:\n",classification_report(y_test, prediction_test))
    skplt.plot_confusion_matrix(y_train,prediction_train,title='Confusion Matrix - Train Data')
    skplt.plot_confusion_matrix(y_test,prediction_test,title='Confusion Matrix - Test Data')
```

XGBoost Classifier

```
In [37]: from xgboost import XGBClassifier

def xgb_Classifier(X_train,X_cv,y_train,y_cv):
    pred_cv = []
    pred_train = []
    depths = [2, 3, 5, 8, 10, 20]
    estimators = [100, 200, 300, 400, 500]
    for i in depths:
        for j in estimators:
            clf = XGBClassifier(n_estimators=j, max_depth=i, scale_pos_weight=1, objective='binary:logistic')
            clf.fit(X_train,y_train)
            prob_train = clf.predict_proba(X_train)[:,:1]
            prob_cv = clf.predict_proba(X_cv)[:,:1]
            auc_score_train = roc_auc_score(y_train,prob_train)
            auc_score_cv = roc_auc_score(y_cv,prob_cv)
            pred_train.append(auc_score_train)
            pred_cv.append(auc_score_cv)
    cmap=sns.light_palette("green")
    # representing heat map for auc score
    print("-"*40, "AUC Score for train data", "-"*40)
    pred_train = np.array(pred_train)
    pred_train = pred_train.reshape(len(depths),len(estimators))
    plt.figure(figsize=(10,5))
    sns.heatmap(pred_train,annot=True, cmap=cmap, fmt=".3f", xticklabels=estimators,yticklabels=depths)
    plt.xlabel('Estimators')
    plt.ylabel('Depths')
    plt.show()
    print("-"*40, "AUC Score for CV data", "-"*40)
    pred_cv = np.array(pred_cv)
    pred_cv = pred_cv.reshape(len(depths),len(estimators))
    plt.figure(figsize=(10,5))
    sns.heatmap(pred_cv, annot=True, cmap=cmap, fmt=".3f", xticklabels=estimators, yticklabels=depths)
    plt.xlabel('Estimators')
    plt.ylabel('Depths')
    plt.show()
```

Testing Model

```
In [41]: import scikitplot.metrics as skplt
def testing1(X_train,y_train,X_test,y_test,optimal_depth,optimal_estimator):
    clf = XGBClassifier(n_estimators = optimal_estimator, max_depth = optimal_depth)
    clf.fit(X_train,y_train)
    prob_train1 = clf.predict_proba(X_train)[:,:1]
    prob_test1 = clf.predict_proba(X_test)[:,:1]

    print("AUC Score for train data",roc_auc_score(y_train,prob_train1))
    print("AUC Score for test data",roc_auc_score(y_test,prob_test1))

    fpr_train, tpr_train, thresholds = roc_curve(y_train,prob_train1)
    fpr_test, tpr_test, thresholds = roc_curve(y_test,prob_test1)

    plt.plot([0, 1], [0, 1], linestyle='--')
    plt.plot(fpr_train, tpr_train, marker='.',color= 'r',label='Train Data')
    plt.plot(fpr_test, tpr_test, marker='.',color = 'b',label='Test Data')
    plt.title("Line Plot of ROC Curve on Train Data and Test Data")
    plt.legend(loc='upper left')
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()
    #plot confusion matrix
    prediction_train=clf.predict(X_train)
    prediction_test=clf.predict(X_test)
    print("macro f1 score for train data :",f1_score(y_train, prediction_train, average = 'macro'))
    print("macro f1 score for test data :",f1_score(y_test, prediction_test, average = 'macro'))
    print("micro f1 score for train data:",f1_score(y_train, prediction_train, average = 'micro'))
    print("micro f1 score for test data:",f1_score(y_test, prediction_test, average = 'micro'))
    print("hamming loss for train data:",hamming_loss(y_train,prediction_train))
    print("hamming loss for test data:",hamming_loss(y_test,prediction_test))
    print("Precision recall report for train data:\n",classification_report(y_train, prediction_train))
    print("Precision recall report for test data:\n",classification_report(y_test, prediction_test))
    skplt.plot_confusion_matrix(y_train,prediction_train,title='Confusion Matrix for Train Data')
    skplt.plot_confusion_matrix(y_test,prediction_test,title='Confusion Matrix for Test Data')
```

Top 20 features:

```
In [23]: from wordcloud import WordCloud

def imp_feature(vectorizer, classifier, n=20):
    features = []
    feature_names = vectorizer.get_feature_names()
    coefs = sorted(zip(classifier.feature_importances_, feature_names))
    top = coefs[:-(n+1):-1]
    print('\033[1m' + "feature_importances\tfeatures" + '\033[0m')
    print("="*35)
    for (coef1, feat1) in top:
        print("%.4f\t\t\t%-15s" % (coef1, feat1))
        features.append(feat1)
    wordcloud = WordCloud(background_color='black', width=1600, height=800).generate(" ".join(features))  #top 20 features in wor
d cloud
    fig = plt.figure(figsize=(30, 20))
    plt.imshow(wordcloud)
    plt.axis('off')
    plt.tight_layout(pad=0)
    #fig.savefig("features.png")
    plt.show()
```

Techniques for vectorization :-

1. Bag of Words (BoW)

```
In [24]: count_vec = CountVectorizer()
BOW_X_train = count_vec.fit_transform(X_tr)
BOW_X_cv = count_vec.transform(X_cv)
BOW_X_test = count_vec.transform(X_test)
```

```
In [25]: #Standardizing data using StandardScaler


sc = StandardScaler(with_mean=False)
BOW_X_train_sc = sc.fit_transform(BOW_X_train)
BOW_X_cv_sc = sc.transform(BOW_X_cv)
BOW_X_test_sc = sc.transform(BOW_X_test)

print("The shape of out text BOW vectorizer ", BOW_X_train_sc.get_shape())
print("CV Data Size: ", BOW_X_cv_sc.shape)
print("Test Data Size: ", BOW_X_test_sc.shape)
```

```
The shape of out text BOW vectorizer (43008, 24467)
CV Data Size: (18433, 24467)
Test Data Size: (26332, 24467)
```

Training RandomForest Classifier:


```
In [31]: RF_Classifier(BOW_X_train_sc,BOW_X_cv_sc,y_tr,y_cv)
```



Testing RandomForest Classifier:


```
In [34]: import scikitplot
testing(BOW_X_train_sc,y_tr,BOW_X_test_sc,y_test,optimal_depth=5,optimal_estimator=400)
```

AUC Score for train data 0.9180368489461701
AUC Score for test data 0.9046712911904043




macro f1 score for train data : 0.7834606832239264
macro f1 score for test data : 0.7693936685354561
micro f1 score for train data: 0.8648856026785714
micro f1 score for test data: 0.8556888956402856
hamming loss for train data: 0.13511439732142858
hamming loss for test data: 0.14431110435971442

Precision recall report for train data:

	precision	recall	f1-score	support
0	0.56	0.77	0.65	6998
1	0.95	0.88	0.92	36010
accuracy			0.86	43008
macro avg	0.76	0.83	0.78	43008
weighted avg	0.89	0.86	0.87	43008


Precision recall report for test data:

	precision	recall	f1-score	support
0	0.54	0.76	0.63	4234
1	0.95	0.87	0.91	22098
accuracy			0.86	26332
macro avg	0.74	0.82	0.77	26332
weighted avg	0.88	0.86	0.87	26332



Training XGBoost Classifier:


```
In [39]: xgb_Classifier(BOW_X_train_sc,BOW_X_cv_sc,y_tr,y_cv)
```



Testing XGBoost Classifier:

```
In [42]: testing1(BOW_X_train_sc,y_tr,BOW_X_test_sc,y_test,optimal_depth=3,optimal_estimator=400)
```

AUC Score for train data 0.9447968987687916
AUC Score for test data 0.9181528695573584





macro f1 score for train data : 0.7830484728416077
macro f1 score for test data : 0.7520082923037057
micro f1 score for train data: 0.9063430059523809
micro f1 score for test data: 0.8947288470302294
hamming loss for train data: 0.09365699404761904
hamming loss for test data: 0.10527115296977062
Precision recall report for train data:

	precision	recall	f1-score	support
0	0.91	0.47	0.62	6998
1	0.91	0.99	0.95	36010
accuracy			0.91	43008
macro avg	0.91	0.73	0.78	43008
weighted avg	0.91	0.91	0.89	43008

Precision recall report for test data:

	precision	recall	f1-score	support
0	0.84	0.42	0.56	4234
1	0.90	0.99	0.94	22098
accuracy			0.89	26332
macro avg	0.87	0.70	0.75	26332
weighted avg	0.89	0.89	0.88	26332




Top 20 Features:

For RandomForest:

```
In [43]: clf = RandomForestClassifier(max_depth =5, n_estimators = 400,class_weight='balanced')
         clf.fit(BOW_X_train_sc,y_tr)
         features = imp_feature(count_vec,clf)
```

feature_importances	features
0.0212	disappoint
0.0197	great
0.0192	horribl
0.0183	bad
0.0174	money
0.0155	love
0.0143	worst
0.0140	aw
0.0139	snack
0.0130	best
0.0124	perfect
0.0124	delici
0.0120	return
0.0117	thought
0.0112	favorit
0.0111	would
0.0109	didn't
0.0109	threw
0.0099	nice
0.0088	receiv



For XGBoost:

```
In [44]: clf = XGBClassifier(max_depth=3, n_estimators = 400, class_weight='balanced')
clf.fit(BOW_X_train_sc, y_tr)
features = imp_feature(count_vec, clf)
```

feature_importances	features
0.0251	disappoint
0.0162	return
0.0148	great
0.0145	horribl
0.0142	love
0.0127	wast
0.0117	money
0.0116	worst
0.0109	terribl
0.0108	aw
0.0097	best
0.0094	bad
0.0093	delici
0.0093	threw
0.0086	refund
0.0083	thought
0.0082	favorit
0.0079	perfect
0.0077	mayb
0.0077	disgust



2. TF-IDF

```
In [45]: tfidf_vec = TfidfVectorizer(ngram_range=(1,2))
tfidf_train = tfidf_vec.fit_transform(X_tr)
tfidf_cv = tfidf_vec.transform(X_cv)
tfidf_test = tfidf_vec.transform(X_test)

print("The type of count vectorizer ", type(tfidf_train))
print("The shape of out text TFIDF vectorizer ", tfidf_train.get_shape())
print("Size of CV dataset:", tfidf_cv.shape)
print("Size of test dataset:", tfidf_test.shape)
```


```
The type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
The shape of out text TFIDF vectorizer (43008, 683623)
Size of CV dataset: (18433, 683623)
Size of test dataset: (26332, 683623)
```

```
In [46]: #Standardizing data using StandardScaler
```

```
sc = StandardScaler(with_mean=False)
tfidf_train_sc = sc.fit_transform(tfidf_train)
tfidf_cv_sc = sc.transform(tfidf_cv)
tfidf_test_sc = sc.transform(tfidf_test)
```

Training RandomForest Classifier


```
In [47]: RF_Classifier(tfidf_train_sc,tfidf_cv_sc,y_tr,y_cv)
```



Testing RandomForest Classifier:

```
In [35]: testing(tfidf_train_sc,y_tr,tfidf_test_sc,y_test,optimal_depth=5,optimal_estimator=400)
```

AUC Score for train data 0.911302287026269
AUC Score for test data 0.8101831770299801




macro f1 score for train data : 0.7911008649423039
macro f1 score for test data : 0.7079217297188739
micro f1 score for train data: 0.8983909970238095
micro f1 score for test data: 0.8580434452377336
hamming loss for train data: 0.10160900297619048
hamming loss for test data: 0.14195655476226646
Precision recall report for train data:

	precision	recall	f1-score	support
0	0.75	0.56	0.64	6998
1	0.92	0.96	0.94	36010
accuracy			0.90	43008
macro avg	0.84	0.76	0.79	43008
weighted avg	0.89	0.90	0.89	43008


Precision recall report for test data:

	precision	recall	f1-score	support
0	0.58	0.44	0.50	4234
1	0.90	0.94	0.92	22098
accuracy			0.86	26332
macro avg	0.74	0.69	0.71	26332
weighted avg	0.85	0.86	0.85	26332



Training XGBoost Classifier


```
In [49]: xgb_Classifier(tfidf_train_sc,tfidf_cv_sc,y_tr,y_cv)
```



Testing XGBoost Classifier:


```
In [50]: testing1(tfidf_train_sc,y_tr,tfidf_test_sc,y_test,optimal_depth=5,optimal_estimator=300)
```

AUC Score for train data 0.9726735388910658
AUC Score for test data 0.9286331845607403





macro f1 score for train data : 0.8504638974360126
macro f1 score for test data : 0.7824111249400486
micro f1 score for train data: 0.9312220982142857
micro f1 score for test data: 0.9042229986328422
hamming loss for train data: 0.06877790178571429
hamming loss for test data: 0.09577700136715783
Precision recall report for train data:

	precision	recall	f1-score	support
0	0.96	0.60	0.74	6998
1	0.93	0.99	0.96	36010
accuracy			0.93	43008
macro avg	0.94	0.80	0.85	43008
weighted avg	0.93	0.93	0.92	43008

Precision recall report for test data:

	precision	recall	f1-score	support
0	0.86	0.49	0.62	4234
1	0.91	0.98	0.95	22098
accuracy			0.90	26332
macro avg	0.88	0.73	0.78	26332
weighted avg	0.90	0.90	0.89	26332



Top 20 Features:

RandomForest

```
In [51]: clf = RandomForestClassifier(max_depth =5, n_estimators = 400, class_weight='balanced')
         clf.fit(tfidf_train_sc, y_tr)
         features = imp_feature(tf_idf_vec, clf)
```


feature_importances	features
0.0174	delici
0.0091	great
0.0089	perfect
0.0073	wast money
0.0070	contact
0.0066	thought
0.0065	return
0.0062	tast great
0.0060	excel
0.0058	guess
0.0058	product
0.0057	high
0.0054	away
0.0053	unfortun
0.0048	terribl
0.0047	money
0.0046	wast
0.0044	aw
0.0044	noth
0.0044	throw



XGBoost

```
In [52]: clf = XGBClassifier(max_depth=5, n_estimators = 300, class_weight='balanced')
clf.fit(tfidf_train_sc, y_tr)
features = imp_feature(tf_idf_vec, clf)
```

feature_importances	features
0.0126	disappoint
0.0124	wast money
0.0112	worst
0.0098	return
0.0089	terribl
0.0084	horribl
0.0080	best
0.0080	money
0.0077	delici
0.0076	perfect
0.0073	favorit
0.0069	great
0.0068	descript
0.0066	aw
0.0066	wast
0.0066	bad
0.0065	high recommend
0.0065	refund
0.0063	easi
0.0063	wont buy



3. Avg-W2V

```
In [53]: i=0
list_sent_train=[]
for sent in X_tr:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    list_sent_train.append(filtered_sentence)
```

```
In [54]: i=0
list_sent_train1=[]
for sent in X_tr:
    filtered_sentence=[]
    sent=sent
    for w in sent.split():
        for cleaned_words in w.split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    list_sent_train1.append(filtered_sentence)
```

```
In [55]: i=0
list_sent_cv=[]
for sent in X_cv:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    list_sent_cv.append(filtered_sentence)
```

```
In [56]: i=0
list_sent_test=[]
for sent in X_test:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    list_sent_test.append(filtered_sentence)
```

```
In [57]: import gensim
w2v_model = gensim.models.Word2Vec(list_sent_train,min_count=5,size=50,workers=4)
w2v_words = list(w2v_model.wv.vocab)
```

```
In [58]: def avg_w2v(list_of_sent):
    sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in list_of_sent: # for each review/sentence
        sent_vec = np.zeros(50) # as word vectors are of zero length
        cnt_words =0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_words:
                vec = w2v_model.wv[word]
                sent_vec += vec
                cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
        sent_vectors.append(sent_vec)
    print(len(sent_vectors))
    print(len(sent_vectors[0]))
    return sent_vectors
```

```
In [59]: train_avgw2v = avg_w2v(list_sent_train)

43008
50
```

```
In [60]: cv_avgw2v = avg_w2v(list_sent_cv)

18433
50
```

```
In [61]: test_avgw2v = avg_w2v(list_sent_test)


26332
50
```

```
In [62]: #Standardizing data using StandardScaler

sc = StandardScaler(with_mean=False)
aw2v_X_train_sc = sc.fit_transform(train_avgw2v)
aw2v_X_cv_sc = sc.transform(cv_avgw2v)
aw2v_X_test_sc = sc.transform(test_avgw2v)
```

Training RandomForest Classifier:


```
In [64]: RF_Classifier(aw2v_X_train_sc,aw2v_X_cv_sc,y_tr,y_cv)
```



Testing RandomForest Classifier:

```
In [65]: testing(aw2v_X_train_sc,y_tr,aw2v_X_test_sc,y_test,optimal_depth=8,optimal_estimator=300)
```

AUC Score for train data 0.925272660915774
AUC Score for test data 0.8792803970700706




macro f1 score for train data : 0.7746311446357439
macro f1 score for test data : 0.7298476718186518
micro f1 score for train data: 0.8481910342261905
micro f1 score for test data: 0.8202567218593346
hamming loss for train data: 0.15180896577380953
hamming loss for test data: 0.17974327814066535
Precision recall report for train data:

	precision	recall	f1-score	support
0	0.52	0.85	0.65	6998
1	0.97	0.85	0.90	36010
accuracy			0.85	43008
macro avg	0.74	0.85	0.77	43008
weighted avg	0.89	0.85	0.86	43008


Precision recall report for test data:

	precision	recall	f1-score	support
0	0.46	0.75	0.57	4234
1	0.95	0.83	0.89	22098
accuracy			0.82	26332
macro avg	0.70	0.79	0.73	26332
weighted avg	0.87	0.82	0.84	26332



Training XGBoost Classifier:

```
In [66]: xgb_Classifier(aw2v_X_train_sc,aw2v_X_cv_sc,y_tr,y_cv)
```




Testing XGBoost Classifier:

```
In [67]: testing1(aw2v_X_train_sc,y_tr,aw2v_X_test_sc,y_test,optimal_depth=3,optimal_estimator=300)
```

AUC Score for train data 0.9317263217744841

AUC Score for test data 0.9024958837331006



macro f1 score for train data : 0.7881317533040173

macro f1 score for test data : 0.7548455368515199

micro f1 score for train data: 0.9007859002976191

micro f1 score for test data: 0.887969011089169

hamming loss for train data: 0.09921409970238096


hamming loss for test data: 0.11203098891083092

Precision recall report for train data:

	precision	recall	f1-score	support
0	0.79	0.53	0.63	6998
1	0.91	0.97	0.94	36010
accuracy			0.90	43008
macro avg	0.85	0.75	0.79	43008
weighted avg	0.89	0.90	0.89	43008

Precision recall report for test data:

	precision	recall	f1-score	support
0	0.74	0.47	0.57	4234
1	0.91	0.97	0.94	22098
accuracy			0.89	26332
macro avg	0.82	0.72	0.75	26332
weighted avg	0.88	0.89	0.88	26332



4. TF_IDF-W2V

```
In [68]: tfidf_vect = TfidfVectorizer()
tfidf_train = tfidf_vect.fit_transform(X_tr)
print("The type of count vectorizer ",type(tfidf_train))
print("The shape of out text TFIDF vectorizer ",tfidf_train.get_shape())
tfidf_cv = tfidf_vect.transform(X_cv)
tfidf_test = tfidf_vect.transform(X_test)
print("CV Data Size: ",tfidf_cv.shape)
print("Test Data Size: ",tfidf_test.shape)
```

The type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>

The shape of out text TFIDF vectorizer (43008, 24467)

CV Data Size: (18433, 24467)

Test Data Size: (26332, 24467)


```
100%|██████████████████████████████████████████████████████████████████████████████| 43008/43008 [16:37<00:00, 43.13it/s]
43008
50
```

```
100%|██████████████████████████████████████████████████████████████| 18433/18433 [06:29<00:00, 40.64it/s]
18433
50
6.487487379533317
```


```
100%|██████████████████████████████████████████████████████████████████████████████| 26332/26332 [10:31<00:00, 55.55it/s]
26332
50
17.020035282533595
```

```
In [73]: #Standardizing data using StandardScaler

sc = StandardScaler(with_mean=False)
tfidf2v_X_train_sc = sc.fit_transform(train_tfidf2v)
tfidf2v_X_cv_sc = sc.transform(cv_tfidf2v)
tfidf2v_X_test_sc = sc.transform(test_tfidf2v)
```

25/30

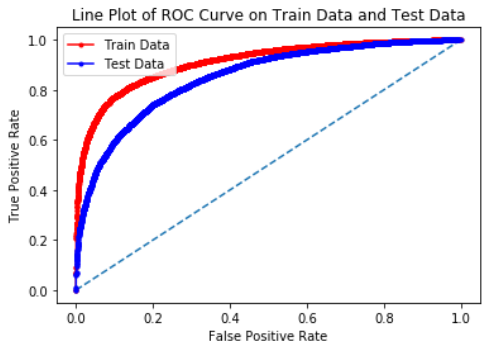
```
In [74]: RF_Classifier(tfidf2v_X_train_sc,tfidf2v_X_cv_sc,y_tr,y_cv)
```



Testing RandomForest Classifier

```
In [75]: testing(tfidf2v_X_train_sc,y_tr,tfidf2v_X_test_sc,y_test,optimal_depth=8,optimal_estimator=400)
```

AUC Score for train data 0.9117938822366751
AUC Score for test data 0.8518323474514458




macro f1 score for train data : 0.7615347925477747
macro f1 score for test data : 0.708712055349226
micro f1 score for train data: 0.8417503720238095
micro f1 score for test data: 0.8103068509797965
hamming loss for train data: 0.15824962797619047
hamming loss for test data: 0.18969314902020357
Precision recall report for train data:

	precision	recall	f1-score	support
0	0.51	0.80	0.62	6998
1	0.96	0.85	0.90	36010
accuracy			0.84	43008
macro avg	0.73	0.83	0.76	43008
weighted avg	0.88	0.84	0.85	43008


Precision recall report for test data:

	precision	recall	f1-score	support
0	0.44	0.68	0.54	4234
1	0.93	0.83	0.88	22098
accuracy			0.81	26332
macro avg	0.69	0.76	0.71	26332
weighted avg	0.85	0.81	0.83	26332



Training XGBoost classifier


```
In [76]: xgb_Classifier(tfidf2v_X_train_sc,tfidf2v_X_cv_sc,y_tr,y_cv)
```



Testing XGBoost Classifier:

```
In [77]: testing1(tfidf2v_X_train_sc,y_tr,tfidf2v_X_test_sc,y_test,optimal_depth=3,optimal_estimator=500)
```

AUC Score for train data 0.9293100305804038
AUC Score for test data 0.8820471765463698





macro f1 score for train data : 0.7748150036016999
macro f1 score for test data : 0.7211724086212047
micro f1 score for train data: 0.8976469494047619
micro f1 score for test data: 0.8779811636032204
hamming loss for train data: 0.1023530505952381
hamming loss for test data: 0.12201883639677959

Precision recall report for train data:

	precision	recall	f1-score	support
0	0.81	0.49	0.61	6998
1	0.91	0.98	0.94	36010
accuracy			0.90	43008
macro avg	0.86	0.73	0.77	43008
weighted avg	0.89	0.90	0.89	43008

Precision recall report for test data:

	precision	recall	f1-score	support
0	0.72	0.40	0.51	4234
1	0.89	0.97	0.93	22098
accuracy			0.88	26332
macro avg	0.81	0.68	0.72	26332
weighted avg	0.87	0.88	0.86	26332



```
In [78]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "Best Hyper Parameter(Depth)", "Best Hyper parameter(n_estimator)", "Test Auc Score"]
x.add_row(["Bow", "Random Forest", 5, 400, 90.4])
x.add_row(["Tf-Idf", "Random Forest", 5, 400, 81.01])
x.add_row(["Avg-W2V", "Random Forest", 8, 300, 87.92])
x.add_row(["TfIdf-W2V", "Random Forest", 8, 400, 85.18])
x.add_row(["Bow", "XGBoost", 3, 400, 91.8])
x.add_row(["Tf-Idf", "XGBoost", 5, 300, 92.86])
x.add_row(["Avg-W2V", "XGBoost", 3, 300, 90.24])
x.add_row(["TfIdf-W2V", "XGBoost", 3, 500, 88.2])
from IPython.display import Markdown, display
def printmd(string):
    display(Markdown(string))
printmd('****Final Conclusion:****')
print(x)
```

Final Conclusion:

Vectorizer	Model	Best Hyper Parameter(Depth)	Best Hyper parameter(n_estimator)	Test Auc Score
Bow	Random Forest	5	400	90.4
Tf-Idf	Random Forest	5	400	81.01
Avg-W2V	Random Forest	8	300	87.92
TfIdf-W2V	Random Forest	8	400	85.18
Bow	XGBoost	3	400	91.8
Tf-Idf	XGBoost	5	300	92.86
Avg-W2V	XGBoost	3	300	90.24
TfIdf-W2V	XGBoost	3	500	88.2